Planning & Design Report

Group 7

Harry Stevenson
Stephen Lynn
David Todd
Liam Joy
Thomas Gak Deluen

February 9, 2017

1 Abstract

This document outlines the design of the system that will meet the requirements outlined in the Requirements Analysis report for the stakeholders, Deutsche Bank. The system will serve to alert the stakeholders to irregularities and anomalies in simulated streams of FTSE 100 stock market data. This solution will be using general statistical analysis methods which can be applied to any field to detect anomalous data, but will be specialised in this case for analytics of trading data. In a real life case, the trading data will be much richer in information, however the same analytical methods will still apply.

2 Group Communication

To communicate between everyone in the group, we will utilise Slack. The coding done shall be pushed to our private repository on GitHub for a centralised code base and its version control. We will use Google Drive and Google Docs for sharing reports, so all team members can work simultaneously on the same document.

As displayed in figure 1, the Gantt chart will be used to ensure that the project runs to time. We have built in buffers such that our group deadline for certain tasks is before the official deadline, such that if tasks overrun for whatever reason, then it should not compromise the whole project. For example, we plan to get the primary features of the coding implemented by the end of week 7, such that it gives us time to implement stretch goals and perform rigorous testing on the system.

The activity network in figure 2 shows both how tasks are assigned to different group members and the dependencies between activities. The tasks have been assigned to group members based on their strengths and the tasks that they have completed earlier on in the project. The dependencies should be such that the majority of tasks should not be waiting upon other tasks, however, when this is unavoidable (for example, writing the final report) then there should be enough material to start the task.
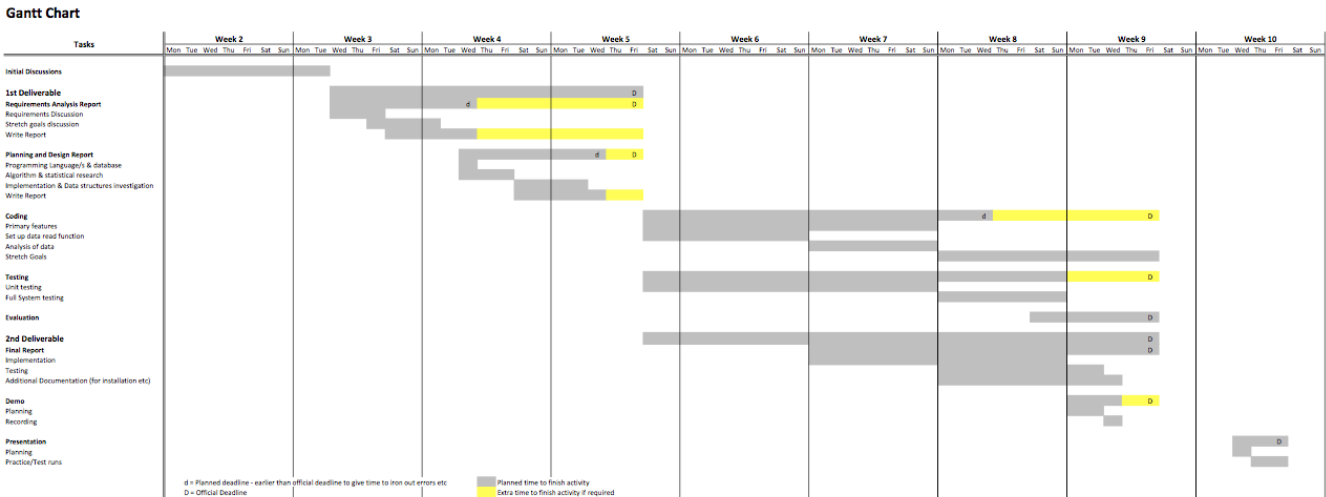

Figure 1: Gantt chart to show when tasks within the project will be completed

3 Software Methodology

We have chosen to adopt an agile methodology to development, with a scrum approach. We decided that the ideal time to have a sprint planning session would be at the beginning of every

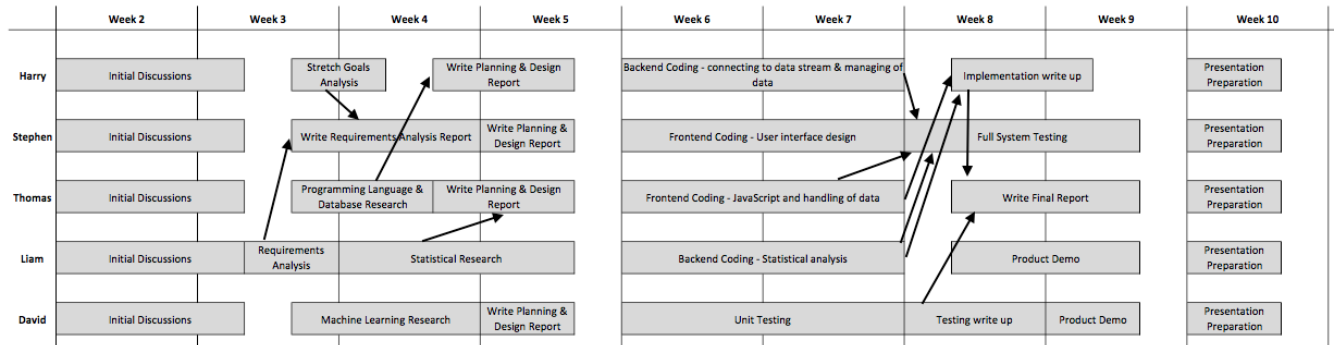|	Week 2	Week 3	Week 4	Week 5	Week 6	Week 7	Week 8	Week 9	Week 10
Harry	Initial Discussions		Stretch Goals Analysis	Write Planning & Design Report	Backend Coding - connecting to data stream & managing of data		Implementation write up		Presentation Preparation
Stephen	Initial Discussions		Write Requirements Analysis Report	Write Planning & Design Report	Frontend Coding - User interface design		Full System Testing		Presentation Preparation
Thomas	Initial Discussions		Programming Language & Database Research	Write Planning & Design Report	Frontend Coding - JavaScript and handling of data		Write Final Report		Presentation Preparation
Liam	Initial Discussions		Requirements Analysis	Statistical Research	Backend Coding - Statistical analysis		Product Demo		Presentation Preparation
David	Initial Discussions		Machine Learning Research	Write Planning & Design Report	Unit Testing		Testing write up	Product Demo	Presentation Preparation

Figure 2: Activity network to visualise how tasks are assigned throughout the group

week, commencing after the planning and design document has been finalised. Given the time constraints of the project, a longer gap between sprint sessions would make an iterative approach more difficult. We chose this methodology over the classical waterfall approach as it permits a greater level of interaction with the customer and therefore allowing the system to better meet the project requirements. In addition, it allows us to first perfect the basic functionality of the system before adding additional features.

4 Backend

4.1 Programming Languages

The backend of the software shall be written in Python. We chose this as it offers a large variety of machine learning libraries to choose from, such as scikit-learn. Python has a very simple and straightforward syntax, which will increase the readability of our code. A large part of the backend of this project is statistical analysis of data. Python balances strong statistical analysis with general purpose usage, a purely statistical language such as R would not suffice.

The system will use a RethinkDB database. This is a NoSQL document-based database, well suited for real-time web applications. It allows the frontend to listen for changes, such as update, delete or insert, therefore producing quick updates to the user when the backend receives real-time market data or sends an alert. This is an important aspect of our system which requires traders to make quick, informed decisions. RethinkDB also has official drivers for Python and JavaScript, which simplifies the process of connection to the database for backend analytics.

4.2 Backend Workflow

As can be seen in the activity diagram in figure 3, each trade is read, either from the live data feed or the static file, then placed into the database. Then, each row is analysed using the above statistical methods. If these methods indicate an irregularity, an alert is triggered on the front end for the user. Otherwise, the current running averages are updated to take the new row into account.

4.3 Anomalous Data Detection

We visualised the sample data that we have been given using histograms to plot the frequency of particular values such as trade sizes or the change in price between trades. We found this data gave a skewness between -1 and 1. Thus, the data can be assumed to be normally distributed and can be modelled using the normal distribution. The following formulae will be used for
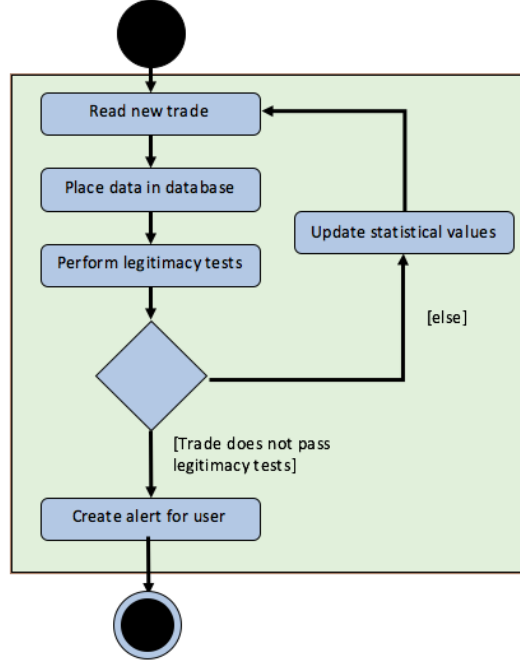
Figure 3: Activity diagram to show the flow of data coming in and out of the system

calculations:

Mean: $\quad \mu = \frac{\sum_{i=1}^{N} x_i}{N}$

Standard Deviation: $\quad \sigma = \sqrt{\frac{1}{N} \sum_{i=1}^{N} (x_i - \mu)^2}$

Expected Frequency Outside Range $= 1$ *in* $\frac{1}{1 - erf(\frac{x}{\sqrt{2}})}$

Where Range $= \mu \pm x\sigma$

Where erf(x) $= \frac{2}{\sqrt{\pi}} \int_{0}^{x} e^{-x^2}$

As stated by the Empirical Rule (68-95-99.7), 99.7% of the data should lie in the range $\mu \pm 3\sigma$. The above formulae show that a value that lies outside the above range should occur only once every approximately 370 values:

$$1 \ in \ \frac{1}{1-erf(\frac{x}{\sqrt{2}})} \implies 1 \ in \ \frac{1}{1-erf(\frac{3}{\sqrt{2}})} = 1 \ in \ \frac{1}{1-0.9973...} = 1 \ in \ 370.37... \approx 1 \ in \ 370$$

This is a good basis to begin to detect anomalous data from.

Other anomalies will involve an increase/decrease in mean values for the size/price of trades. This could indicate fraudulent practices such as pump and dumps or bear raids. Volume spikes could also be picked up by analysing a change in size of trades for particular company stocks. Using statistical analysis methods like the ones described above could flag a variety of anomalies in the data.

The histogram in figure 4, constructed from the supplied data, where the price of a trade of a single stock is compared to the previous trade price of that stock to find the difference between trades. The data has a skewness of 0.0113 and so can be assumed as normally distributed. This is an example of how you can detect anomalies that occur at the very edges of the graph as they are least likely to occur naturally.

As a stretch goal for our project, we will attempt to categorise these different anomalies so that our program can learn from the data and alert the user to a possible reason for the anomaly. This could include distinguishing between a volume spike and an intended artificial

3

Figure 4: Histogram showing the price change against frequency and the fit of the normal distribution

inflation of a stock seen in pump and dumps, both would raise the price of a stock and therefore may appear at first to be similar anomalies.

4.4 Machine Learning Algorithm

For our predictions and some of our analysis, our system will use linear regression to find a line of best fit between our data points. The line will be of the form $y = mx + c$ with m and c being calculated like so:

$$m = \frac{\bar{x}.\bar{y} - \bar{xy}}{(\bar{x})^2 - \bar{x^2}} \qquad\qquad c = \bar{y} - m\bar{x}$$

Where $\bar{x}$ is the mean of all x values of our data points, and $\bar{y}$ is the mean of all y values of the data points. X values will be representing time, and Y values will represent the stock characteristics such as prices and volume. Once we have this line, we will be able to extend it beyond our current data to enable us to predict future Y values, within a range. Since we will be using unpredictable data that can change rapidly, we will vary the size of the set of data points, thus the predictions for future prices will be in the form of a range of values. We will allow the user to define how many trades are used for this analysis and prediction, in turn defining the base size of our sets.

After we have our best fit line, well then be able to calculate how volatile or stable each stock is. Well do this by comparing our actual values to our predicted ones to find the average error e. This formula is shown below:

$$e = \frac{\sum_{i=1}^{n} |y_i - (mx_i + c)|}{n}$$

In our program, we will compare each stocks average error to determine its volatility.

4.5 Class Interaction

Our design features many elements that must work together, the UML class diagram in figure 5 illustrates the relationships between different entities in the system.

The system will have many trades, each for which, two traders with unique ID's will be associated, the buyer and seller. The trade itself will have the attributes as taken from the data stream and have associated methods to compare the price or size or the new trade compared to the statistical data. If the trade does not pass the test against the statistical data, then that one trade will invoke one alert consisting of the stock that was being traded along with the type of anomaly that the offending trade may be categorised as, this alert may then be either approved or rejected. For each single stock there will exist multiple trades and multiple alerts,

4

here each stock will have attributes such as the average price and a trend line which will be used to perform the statistical analysis as well as making predictions based on previous, historical data. A stock may have multiple days of this historical data consisting of the open and close price for the relevant days.
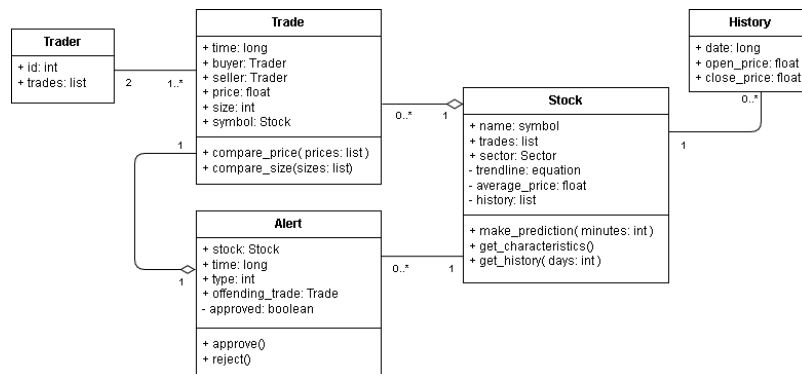

Figure 5: UML class diagram to show how the different entities in the system will interact with each other

5 Frontend

5.1 Programming Languages

We decided that a web application for the front end would offer the greatest level of usability, providing a clear and concise interface to detail the data analysis. In addition to this, it means that it will be compatible with a large range of devices.

The backend and frontend are only connected through the database, leaving room for future changes. If the stakeholders require a different frontend, altering it will be possible without alterations to the backend. Furthermore, a web portal is easily accessible for a large team of users, each with different roles. The specification requires that the system must be suitable for non technical users, and a web-based interface offers an intuitive, user friendly environment to suit such users.

We will build the front-end using JavaScript, to allow for rapid prototyping and real-time updates to the site, alongside it being both fast to execute and fast to write. Furthermore, JavaScript has many prewritten environments, frameworks and applications, which allow us to quickly implement our system with widely used applications.

5.2 Frontend Architecture

The frontend runs in a simple NodeJS app. This uses ExpressJS to handle the different aspects of the server such as routing and static assets. This couples with a ready to use real-time web framework Horizon. Horizon allows the client to connect to the server through web sockets to listen for changes in the database.

We will build the client using ReactJS, a fast and reliable user-interface framework that uses a virtual DOM. This removes the need for DOM manipulations: when the state changes, the required components of the client are redrawn, allowing rapid changes in the database to be quickly propagated to the view. ReactJS is component-based which gives the ability to write components for each part of the client, make testing more in-depth and component specific, allowing bugs to be quickly found and changes in the code to have fewer side-effects as each component works independently, with most components being stateless and presentational. We will write the JavaScript in high level ES7 code and compile it using the Babel compiler. This

5

will maximise the number of supported browsers. It also allows us to test our code every time we compile it.

5.3 Customer Interaction

Our site will consist of three main pages; the homepage, where the user can immediately access their alerts and upload static files. The alerts page, where the system holds every alert for review, along with previous saved alerts. The stocks page, where the user can analyse the performance of all the FTSE 100 stocks. The page includes statistics such as volatility and average stock price alongside short-term stock price predictions. This is illustrated in figure 6. Mock-up design images for each web page in our system are included in the appendix.
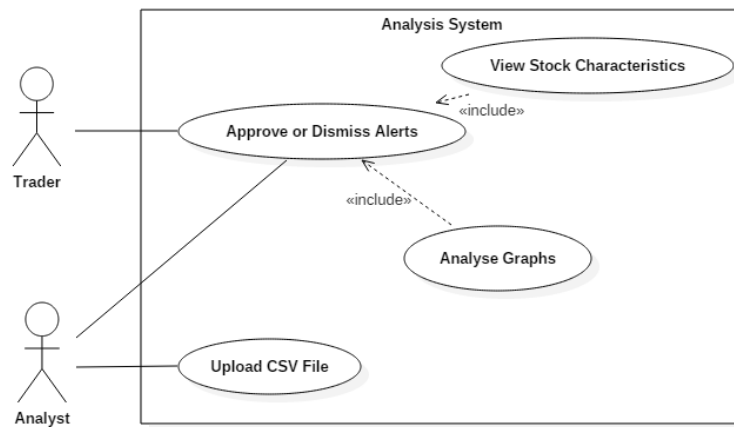

Figure 6: UML use case diagram

6 Architectural Pattern

For this project, we will use the Pipe and Filter architectural pattern. In this case, data is input into the backend, processed, then written to the database. The frontend then reads the data from the database, in addition to receiving data in the form of a CSV file from the customer, which it passes to the backend for analysis. This is illustrated in figure 7.
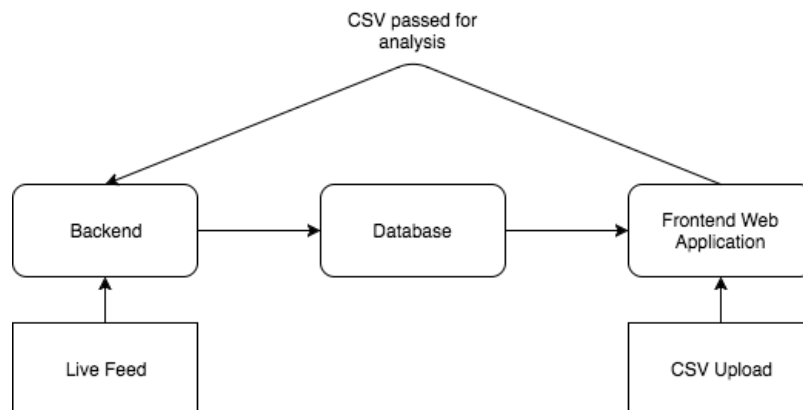

Figure 7: UML architectural pattern diagram showing how data will progress through the system

To keep maintain high uptime and prevent the front end from hanging, once a user uploads a file to the front-end, the frontend writes it to the disk and a forks a background task to analyse the file. Once the backend launches the task, the front-end server lets the user know the file is currently undergoing analysis. The system notifies the user in real-time of the files analysis.

7   Extensibility

The nature of the specification lends the system to further extension. We designed the software so that further statistical analysis can be done on the data if it is required. The agile approach to development will help us to construct the software in this way.

The largest problem faced by the system with regard to extensibility is the storage space required to store all of the trade data. In a CSV file format, one day of trade data is approximately 100MB, around 800,000 rows. This means that when the system is run, it must be run on a machine with a large memory capacity and a high read/write speed.

This system is intended as a model for finding anomalous data. In a real-life case, the data received from the stock market will be much richer in information, however we can use the same fundamental principles implemented in this solution to analyse richer data, and identify different anomalies.

8   Robustness

The system only has two forms of input, a static CSV file, and a live data feed. The live data feed will always be in the same format, so we do not need to validate it. If the CSV file inputted is not correctly formatted, the program will not analyse it, and display a formatting error to the user. Changes to the extraction of the data are simple, as the format used for this project has been simplified for this task, so this can be easily extended to handle more verbose data.

The system will be robust with regard to the live feed of data. As the connection is open, we will check the connections activity frequently, and upon disconnection, it will reconnect as soon as possible, so as little data is lost as possible.

9   Reliability

We will measure the reliability of the system through the consistency of alerts. The alerts will appear within 5 seconds of the analysis of the offending trade, or final trade in a set of trades. We will multithread the backend Python so that we can analyse many trades at once, ensuring that the analytics will be able to keep up with the live feed of data. Furthermore, the RethinkDB database ensures high-speed communication between the frontend and backend. It is able to alert the front end immediately after we detect anomalous activity. In this project, the underlying DB acts as both storage, and a bidirectional messenger between both the frontend and the backend.

10   Correctness

After we build each iteration of the software, we will analyse that iteration to see what requirements it meets, and whether it should have met any other requirements. This is in addition to the unit tests that we will perform at each stage throughout development. This allows us to shape future iterations in a customer driven manner, ensuring that we meet every requirement.

11 Compatibility & Portability

The system will be compatible with a wide range of platforms, including mobile devices, as we are using a web application. We chose this due to the requirements stating that the system should be developed for a wide range of non-technical users. The higher the number of users, the more practical a web application becomes, as long installations are not necessary. The backend of the system should run on a single server. Slow disk access can cause a bottleneck in the speed of the system.

12 Modularity & Reuse

We designed both the backend and frontend in a modular way, so that they only interact through the database. This means that both the frontend and the backend are entirely changeable, provided they write to the database in the same format as the current system.

13 Security

There are no plans in the system to implement a user access system, as this is out of scope of the projects requirements. This means that whoever is able to access the web portal has access to all the analysis of the live feeds data. However, this software only serves as a model to provide data analysis. The stakeholders can implement security aspects of the software when they receive the final product. For example, the system could only be accessible through the stakeholders local network or virtual private network. Access to the system could be incorporated into an SSO (Single Sign On) system that may already be used by the stakeholders for other means. Security considerations that arise from the creation of the software are out of the scope of this project.

14 Risk Analysis

In order to analyse the potential risks to the progress of the project, we have drawn up a risk register. This displays the identified potential events that could cause issues in the project, we have given a probability based on the event occurring and the impact it would have on the project. Where possible we have mitigated against the event from occurring, but in addition, we have also provided a contingency plan, such that there would be less uncertainty if such an event were to occur. As a consequence this reduces the possible impact that the event would have upon the project.

Title	Probability	Impact on	Mitigation	Contingency	Mitigated Impact
Staff Illness	High	Medium	N/A	Delegate role to other team member with lowest current workload	Low
Data Loss of Documentation	Low	High	Cloud based backups through Google Drive	Restore files form Google Drive	Mitigated
Hardware Fault	Very Low	High	N/A	Restore files from that developer's machine to new machine	Mitigated
Team member shirking responsibility	Very Low	High	Reviews of progress in sprint planning sessions	Meet with module tutor to develop arrangement	Low
Competition from other teams	Certain	Medium	Produce highest quality software by deadline	N/A	Low
Change in Requirements	Very Low	Medium	N/A	Analyse new requirements, implement in next agile iteration	Low
Delay in Deliverables	Very Low	Very High	Gantt chart used for planning schedule, progress checks	Meet with customer to rearrage deadlines	High
Staff Turnover	Very Low	High	N/A	Reassign that team member's tasks to other team members	Medium

Figure 8: Risk register

15 Testing Plan

We will test our Python code using built-in unit tests. Unit tests allow early spotting of logical errors as well as checking whether the modules function as intended.

We will also test our JavaScript code using different testing methods as React components display as HTML. Because the JavaScript code is compiled, any syntactic errors are quickly spotted. Therefore, additional tests serve to verify the codes semantic behaviour. After each sprint planning session, we will run manually created unit tests to analyse manually created data designed trigger certain functionalities of the software.

As we are developing the frontend and backend as separate entities, the system requires integration testing to make sure that they can both communicate data between them.

References

Facebook *React*
https://facebook.github.io/react/
Date Accessed: 09/02/2017

Nodejs *Expressjs*
http://expressjs.com/
Date Accessed: 09/02/2017

Nodejs *Docs*
https://nodejs.org/en/docs/
Date Accessed: 09/02/2017

RethinkDB *RethinkDB Horizon Docs*
http://horizon.io/docs/
Date Accessed: 09/02/2017

RethinkDB *RethinkDB*
https://www.rethinkdb.com/
Date Accessed: 09/02/2017

Weisstein, Eric *Erf*
http://mathworld.wolfram.com/Erf.html
Date Accessed: 09/02/2017

Niles, Robert *Standard Deviation*
http://www.robertniles.com/stats/stdev.shtml
Date Accessed: 09/02/2017

Andale *Statistics How To*
http://www.statisticshowto.com/empirical-rule-2/
Date Accessed: 09/02/2017

Bas, Ryan *React Reconciliation*
https://dev.to/ryanbas21/react-reconciliation
Date Accessed: 09/02/2017

Weisstein, Eric *Linear Regression*
http://mathworld.wolfram.com/LinearRegression.html
Date Accessed: 09/02/2017

A User Interface Designs

A.1 Homepage

Figure 9: Mock-up design of home page with annotations

Figure 10: Difference to home page when there are alerts that have yet to be actioned

A.2 Alerts Page

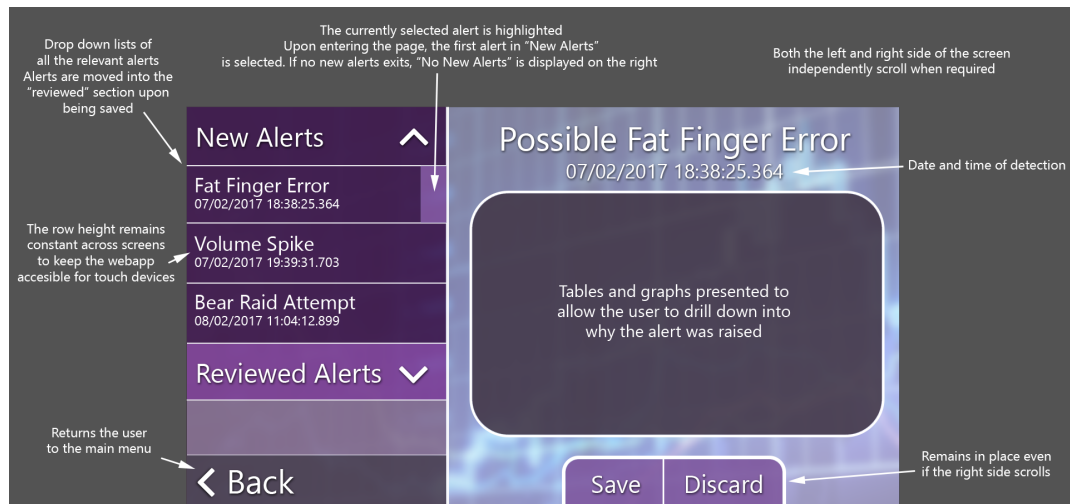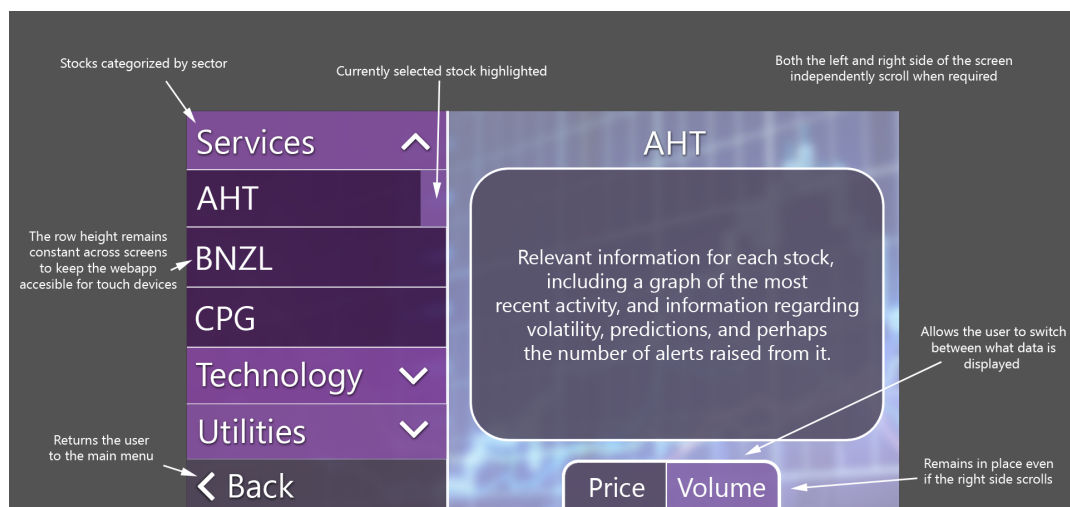Figure 11: Mock-up design of alerts page with annotations

A.3 Stocks Page

Figure 12: Mock-up design of stocks page with annotations